

```

import numpy as np
import random
import warnings
from numpy.random import RandomState

class MAPEstimator():
    """
    Maximum A-Posteriori Estimator for musk probabilities

    Attributes
    -----
    w_D : D-dimensional vector of reals
        Defines weight vector
    prior : string
        Defines prior that will be used options: ('trivial','sas')
        'sas' stands for 'spike-and-slab'
    alpha : float, must be greater than 0
        Defines precision parameter of the multivariate Gaussian prior for the weight vector or the "spike" mul
    beta : float, must be greater than 0 (optional)
        Defines precision parameter of the "slab" Gaussian on the GMM prior for the weight vector
    max_iter: integer
        Defines max number of iterations model can take to converge on fit()
    step_size: float must be greater than 0
        The step size of the gradient descent algorithm
    max_iter: int greater than 0
        Maximum number of iterations that the gradient descent algorithm will take

    Examples
    # TODO : Update this example or delete it
    -----
    >>> word_list = ['dinosaur', 'trex', 'dinosaur', 'stegosaurus']
    >>> mapEst = MAPEstimator(Vocabulary(word_list), alpha=2.0)
    >>> mapEst.fit(word_list)
    >>> np.allclose(mapEst.predict_proba('dinosaur'), 3.0/7.0)
    True

    >>> mapEst.predict_proba('never_seen-before')
    Traceback (most recent call last):
    ...
    KeyError: 'Word never_seen-before not in the vocabulary'
    """

    def __init__(self, w_D, prior='trivial', step_size_type = 'universal', alpha=1.0, beta=None, max_iter=30000, to

        #TODO0 decide on max_iter
        self.w_D = w_D
        self.c = 0
        self.alpha = float(alpha)
        self.max_iter = max_iter
        self.tol = tol
        self.step_size = step_size
        self.step_size_type = step_size_type
        self.iteration_count = iteration_count
        self.prior = prior
        self.loss_array = []
        if prior != 'trivial':
            self.beta = beta

    def fit(self, train_X, train_y):
        """
        Fit this estimator to provided training data with first order stochastic gradient descent

        Args
        ----
        train_X : NxD array
            Each entry is a D-dimensional vector representing a training example
        train_y : Nx1 array
            Each entry is corresponding output class for training data

        Returns
        -----
        None. Internal attributes updated.

        Post Condition

```

```

-----
Attributes will updated based on provided word list
* The 1D array count_V is set to the count of each vocabulary word
* The integer total_count is set to the total length of the word list
'''
self.iteration_count = 0

example_num = 0
num_examples = len(train_X)

self.loss_array = []

if self.prior == 'trivial':

    count_10 = 0
    loss_10 = []
    L_avg = np.inf
    diff_L = np.inf
    L_avg_prev = np.inf

#         while(self.iteration_count <= self.max_iter) and (diff_L > self.tol):
#         # TODO : spit out a warning if the max_iter is reached
#         #print('iteration: %i' % self.iteration_count)

#         h_x = np.matmul(self.w_D.transpose(), train_X) + self.c
#         #print('h_x: ' + str(h_x))
#         # print('w_D: ' + str(self.w_D))
#         # print('c: ' + str(self.c))
#         sig = 1 / (1 + np.exp(-h_x))
#         # print('sig: ' + str(sig))
#         # print('train_y: ' + str(train_y[example_num]))
#         L = np.dot(train_y, np.log(sig)) + np.dot((1-train_y), np.log(1-sig))
#         print('L:' + str(L))
#         loss_10.append(L)
#         count_10 += 1

#         if count_10 == 1000:
#         diff_L = np.abs(L_avg_prev - L)
#         L_avg_prev = L

#         self.w_D = self.w_D - self.step_size * (np.matmul(train_X.transpose(),(sig - train_y)) + self.alp
#         self.c = self.c - self.step_size * (sig - train_y)

#         if example_num >= num_examples:
#         example_num = 0
#         self.iteration_count += 1

#         if self.iteration_count == self.max_iter:
#         warnings.warn("Maximum iterations reached")

prng = RandomState(136)

while(self.iteration_count <= self.max_iter) and (diff_L > self.tol):
# TODO : spit out a warning if the max_iter is reached
# print('iteration: %i' % self.iteration_count)
example_num = prng.randint(0, num_examples-1)
h_x = np.dot(self.w_D, train_X[example_num]) + self.c
# print('h_x: ' + str(h_x))
# print('w_D: ' + str(self.w_D))
# print('c: ' + str(self.c))
sig = 1 / (1 + np.exp(-h_x))
# print('sig: ' + str(sig))
# print('train_y: ' + str(train_y[example_num]))
L = train_y[example_num] * np.log(sig) + (1-train_y[example_num]) * np.log(1-sig)
# print('L:' + str(L))
loss_10.append(L)
count_10 += 1

if count_10 == 100000:
    L_avg = np.mean(loss_10)
    # print(L_avg)
    self.loss_array.append(L_avg)
    diff_L = np.abs(L_avg_prev - L_avg)
    count_10 = 0
    loss_10 = []
    L_avg_prev = L_avg

```

```

        if self.step_size_type == 'universal':
            self.w_D = self.w_D - self.step_size * ((sig - train_y[example_num]) * train_X[example_num] + s
            self.c = self.c - self.step_size * (sig - train_y[example_num])
        else:
            self.w_D = self.w_D - (self.step_size + train_y[example_num]*0.1)* ((sig - train_y[example_num]
            self.c = self.c - (self.step_size + train_y[example_num]*0.1) * (sig - train_y[example_num])

#         example_num += 1
#         if example_num >= num_examples:
#             example_num = 0
        self.iteration_count += 1

        if self.iteration_count == self.max_iter:
            warnings.warn("Maximum iterations reached")

    #else:
    #    TODO: Code up the SGD for the sas prior

def predict_proba(self, test_X):
    """
    Predict probability of a given set of feature vectors under this model

    Args
    ----
    test_X : NxD vector
        Examples for which probability will be predicted

    Returns
    ----
    proba : float between 0 and 1

    Throws
    ----
    ValueError if hyperparameters do not allow MAP estimation
    KeyError if the provided word is not in the vocabulary
    """

    lin_preds = np.matmul(test_X, self.w_D)

    sig_preds = 1 / (1 + np.exp(-lin_preds))

    return sig_preds

def score(self, test_X, test_y):
    """ Compute the average log probability of words in provided list

    Args
    ----
    test_X : NxD array
        Each entry is a D-dimensional vector representing a test example
    test_y : Nx1 array
        Each entry is corresponding output class for test data

    Returns
    ----
    avg_log_proba : float between (-np.inf, 0.0)
    """
    correct_count = 0
    num_examples = len(test_X)

    pred_proba = self.predict_proba(test_X)
    pred_class = pred_proba > 0.5
    is_correct = pred_class == test_y

    return np.sum(is_correct) / num_examples

```

test_gd

April 17, 2022

```
[1]: ## Abstractions
import numpy as np
import pandas as pd

## Plotting
from matplotlib import pyplot as plt
plt.style.use('seaborn')
import seaborn as sns
import pylab as pl

## Scalers
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler

## Models
from sklearn.linear_model import LogisticRegression

## Model Selection
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold
from sklearn.model_selection import StratifiedKFold

## Timing
import time

# Model
from MAPEstimator import MAPEstimator
```

Import Data

```
[2]: headers = ['molecule_name', 'conformation_name']
for i in range(1, 167):
    name = 'f%i' % i
    headers.append(name)
headers.append('class')
```

```
[3]: # headers = pd.read_csv('clean2.info')
df = pd.read_csv('src/clean2.data')
```

```
df.columns = headers
```

```
[4]: X = np.asarray(df.iloc[:,2:-1])  
y = np.asarray(df.iloc[:, -1])
```

Standard Scaler

```
[5]: X_std = StandardScaler().fit_transform(X)
```

Train and Test Model

```
[6]: clf = MAPEstimator(w_D = np.zeros(X.shape[1]), step_size=0.1, alpha=0.1,  
↳max_iter = 10000000)  
clf.fit(X_std,y)  
predict_y = clf.predict_proba(X_std)
```

```
[7]: score = clf.score(X_std, y)  
score
```

```
[7]: 0.6296801576474155
```

```
[8]: (np.sum(y==0)) / len(y)
```


```
[8]: 0.8459906017886918
```

```
[10]: clf.iteration_count
```

```
[10]: 1700000
```

```
[11]: plt.plot(clf.loss_array)
```

```
[11]: [<matplotlib.lines.Line2D at 0x7fd71c99fee0>]
```



```
[17]: clf = MAPEstimator(w_D = np.zeros(X.shape[1]), step_size=0.1, alpha=0.1,
    ↪max_iter = 10000000, step_size_type = 'differential')
    clf.fit(X_std,y)
    predict_y = clf.predict_proba(X_std)
```

```
[18]: score = clf.score(X_std, y)
    score
```

```
[18]: 0.6396847051690162
```

```
[19]: (np.sum(y==0)) / len(y)
```


```
[19]: 0.8459906017886918
```

```
[20]: clf.iteration_count
```

```
[20]: 300000
```

```
[21]: plt.plot(clf.loss_array)
```

```
[21]: [<matplotlib.lines.Line2D at 0x7fd71cb516a0>]
```



Model Evaluation

```
[6]: kf = KFold(n_splits=10, shuffle = True, random_state = 136)
kf.get_n_splits(X_std, y)
```

```
[6]: 10
```

```
[7]: iteration_counts = []
test_scores = []

for train_index, test_index in kf.split(X_std, y):
    X_train, X_test = X_std[train_index], X_std[test_index]
    y_train, y_test = y[train_index], y[test_index]
    clf = MAPEstimator(w_D = np.zeros(X.shape[1]), step_size=0.1, alpha=0.1,
↳max_iter = 10000000)
    clf.fit(X_train, y_train)
    iteration_counts.append(clf.iteration_count)
    score = clf.score(X_test, y_test)
    test_scores.append(score)
```

```
[8]: test_scores
```

```
[8]: [0.6181818181818182,  
      0.5409090909090909,  
      0.6151515151515151,  
      0.6439393939393939,  
      0.5772727272727273,  
      0.6,  
      0.5590909090909091,  
      0.575113808801214,  
      0.6236722306525038,  
      0.6585735963581184]
```

```
[9]: iteration_counts
```

```
[9]: [470000,  
      2986000,  
      155000,  
      150000,  
      1942000,  
      4821000,  
      352000,  
      507000,  
      1327000,  
      260000]
```

```
[10]: basic_test_scores = test_scores
```

```
[11]: basic_iteration_counts = iteration_counts
```

```
[12]: np.mean(basic_test_scores)
```

```
[12]: 0.6011905090357291
```

```
[13]: np.mean(basic_iteration_counts)
```

```
[13]: 1297000.0
```

```
[14]: kf = KFold(n_splits=10, shuffle = True, random_state = 136)  
kf.get_n_splits(X_std, y)
```

```
[14]: 10
```

```
[15]: iteration_counts = []  
test_scores = []  
  
for train_index, test_index in kf.split(X_std, y):  
    X_train, X_test = X_std[train_index], X_std[test_index]  
    y_train, y_test = y[train_index], y[test_index]
```



```

    clf = MAPEstimator(w_D = np.zeros(X.shape[1]), step_size=0.1, alpha=0.1,
↳max_iter = 10000000, step_size_type = 'differential')
    clf.fit(X_train,y_train)
    iteration_counts.append(clf.iteration_count)
    score = clf.score(X_test, y_test)
    test_scores.append(score)

```

/Users/nathanieldavis/Documents/tufts/2022spring/cs136/project/cs136_final_project/Checkpoint 2/MAPEstimator.py:141: RuntimeWarning: divide by zero encountered in log

```

    L = train_y[example_num] * np.log(sig) + (1-train_y[example_num]) *
np.log(1-sig)

```

/Users/nathanieldavis/Documents/tufts/2022spring/cs136/project/cs136_final_project/Checkpoint 2/MAPEstimator.py:141: RuntimeWarning: invalid value encountered in double_scalars

```

    L = train_y[example_num] * np.log(sig) + (1-train_y[example_num]) *
np.log(1-sig)

```

[16]: test_scores

```

[16]: [0.6393939393939394,
      0.6287878787878788,
      0.6272727272727273,
      0.696969696969697,
      0.6681818181818182,
      0.6106060606060606,
      0.5015151515151515,
      0.5386949924127465,
      0.5766312594840668,
      0.6176024279210925]

```

[17]: iteration_counts

```

[17]: [273000,
      889000,
      1997000,
      272000,
      301000,
      411000,
      1267000,
      2137000,
      434000,
      751000]

```

[18]: differential_test_scores = test_scores

[19]: differential_iteration_counts = iteration_counts

```
[20]: np.mean(differential_test_scores)
```

```
[20]: 0.6105655952545179
```

```
[21]: np.mean(differential_iteration_counts)
```

```
[21]: 873200.0
```

```
[22]: plot_test_scores = np.array([np.mean(basic_test_scores), np.  
    ↪ mean(differential_test_scores)])
```

```
[23]: plot_iteration_counts = np.array([np.mean(basic_iteration_counts), np.  
    ↪ mean(differential_iteration_counts)])
```


```
[24]: plot_test_scores = np.array([basic_test_scores, differential_test_scores])
```

```
[25]: plot_iteration_counts = np.array([basic_iteration_counts,   
    ↪ differential_iteration_counts])
```

```
[26]: fig, (ax1, ax2) = plt.subplots(1, 2)  
fig.suptitle('Horizontally stacked subplots')  
ax1.plot(['basic', 'differential'], plot_test_scores)  
ax2.plot(['basic', 'differential'], plot_iteration_counts)
```


```
[26]: [<matplotlib.lines.Line2D at 0x7f9bf0f40c10>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f40bb0>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f40d60>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f40e80>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f40fa0>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f4e100>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f40c40>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f4e220>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f4e430>,  
    <matplotlib.lines.Line2D at 0x7f9bf0f4e550>]
```

Horizontally stacked subplots




```
[35]: plt.boxplot(np.transpose(plot_test_scores), labels = ['static', 'dynamic'])
plt.title('Held-out Accuracy Rates by Weight Update Scheme')
plt.ylabel('Accuracy Rate on Held-out Data')
```

```
[35]: Text(0, 0.5, 'Accuracy Rate on Held-out Data')
```



```
[36]: plt.boxplot(np.transpose(plot_iteration_counts), labels = ['static', 'dynamic'])  
plt.ylabel('Number of Iterations until Convergence (Millions)')
```

```
[36]: Text(0, 0.5, 'Number of Iterations until Convergence (Millions)')
```




```
[51]: plot_test_scores
```

```
[51]: array([[0.61818182, 0.63787879, 0.61515152, 0.64393939, 0.59848485,
            0.5530303 , 0.55909091, 0.57511381, 0.53717754, 0.6585736 ],
            [0.63939394, 0.62878788, 0.63484848, 0.6969697 , 0.66818182,
            0.61060606, 0.65151515, 0.61001517, 0.57663126, 0.61760243]])
```

```
[39]: plt.plot(clf.loss_array)
```

```
[39]: [<matplotlib.lines.Line2D at 0x7f9bf15eed60>]
```



[]: